

Performance Report

Sort on Hadoop/Spark

**Pavankumar Shetty
CWID-A20354961
CS-553 Spring 2016
Main Campus**

INTRODUCTION

This programming assignment covers the Sort application implemented in 3 different ways: **Java (Shared Memory), Hadoop, and Spark.**

The sorting application needed to read a large file (**10GB on single node and 100GB on 16 nodes**) and sort it in place.

As part of this experiment, created the datasets, 10GB dataset and 100GB dataset using Gensort when instances were in running state rather than storing the dataset on Amazon S3. All the experiments were performed on **c3.large** instances.

(i) Shared memory sort:

Sort is an application which sorts a file-resident dataset; As per the requirement, our sort function should be able to sort datasets that are larger than memory of our instances. Sort was performed in multi-threaded environment on 1 virtual node. Time to execute the Sort application on the 10GB dataset produced with the Gensort on 1 node. The multi-threaded experiment was conducted by varying the number of threads from 1 to 8 to check the performance.

Implementation - Data Structures and Algorithms used:

External Merge Sort

It is used as the data being sorted doesn't fit into the main memory of a c3.large instance RAM size. So they should be put locally to the disk, which has slower access rate. Hence breaking down the main file as intermediate files to fit in the memory sounds to be a better approach.

So I used a sort and merge-sort strategy.

Sorting phase involves reading the chunks of data, sort it, and write it out to a temporary file.

In the next phase, some small chunk of data is taken from every temporary files and put it in the memory and sorted and the least key is put to the final output file. This is repeated till every chunk of read from the temporary files.

Execution and Environment setup:

```
ubuntu@ip-172-31-2-233:~$  
ubuntu@ip-172-31-2-233:~$ echo "Shared Memory Sort of 10GB data"
```

Shared Memory Sort of 10GB data

*** Place the java file "SharedMemoryOneThread.java" in the instance, which has java JRE and JDK installed.

```
*** Compile the java file using javac command,  
ubuntu@ip-172-31-2-233:~$ javac SharedMemoryOneThread.java  
ubuntu@ip-172-31-2-233:~$
```

*** Generate the 10GB unsorted data using the gensort command

```
ubuntu@ip-172-31-2-233:~$ ./64/gensort -a 100000000 10GBUnsorted
```

*** Execute the java file with 2 parameters input and output filenames.


```
ubuntu@ip-172-31-2-233:~$ java SharedMemoryOneThread  
10GBUnsorted 10GBSorted
```

*** Convert the output file into dos format to run the valsor


```
ubuntu@ip-172-31-2-233:~$ unix2dos 10GBSorted  
unix2dos: converting file 10GBSorted to DOS format ...
```


No of threads	Execution Time	Serial time Execution	SpeedUp Factor
1	442	442	1
2	423	442	1.044917258
4	435	442	1.016091954
8	439	442	1.006833713

Execution Chart:



Speedup Line Chart:



(ii) Hadoop sort:

The Apache Hadoop software library is a framework that allows for the distributed processing of large data sets across clusters of computers. It is designed to scale up from single servers to thousands of machines (and in our experiment, we have formed a cluster of 16 nodes), each offering local computation and storage.

A MapReduce job usually splits the input data-set into independent splits which are processed by the map tasks in a completely parallel manner.

The framework sorts the outputs of the maps and are fed as input to the reduce tasks. Every job is stored in a Hadoop Distributed File-System (combines together the file systems on all data nodes to make them into one large file system). The Hadoop framework takes care of scheduling tasks, monitoring them and re-executes the failed tasks using JobTracker and TaskTracker.

1. Single Node Sorting

MapReduce is done on the dataset on single node.

Instance type: c3.large

Hadoop version: Hadoop 2.7.2

Execution and Environment setup:

Open up the c3.large instance with java installed AMI.

Send the files, "HadoopPshetty4New.pem" file, 64 folder(which has Gensort and Valsort) and connect to it from local system and follow the following steps to setup and run the sort program on hadoop.

```
// add the pem file to the instance
```

```
eval `ssh-agent -s`
```

```
chmod 600 HadoopPshetty4New.pem
```

ssh-add HadoopPshetty4New.pem

// Create a jar for your mapper reducer sorting java file and place it in bin of hadoop folder.

// before creating jar, add this below path in ./bashrc

export HADOOP_CLASSPATH=/usr/lib/jvm/java-1.7.0-openjdk-amd64/lib/tools.jar

vim HadoopSortSingleNode.java

./hadoop com.sun.tools.javac.Main HadoopSortSingleNode.java

jar cf HadoopSortSingleNode.jar HadoopSortSingleNode*.class

// Change the configuration files in **~/hadoop-2.7.2/etc/hadoop\$**

1. vi core-site.xml

// Add the below configuration in this file.

// **fs.defaultFS** - NameNode URL and port should be specified here

// **hadoop.tmp.dir** - HDFS will be configured in this path.

<configuration>

<property>

<name>fs.defaultFS</name>

<value>hdfs://ec2-54-175-61-81.compute-1.amazonaws.com:9000</value>

</property>

<property>

<name>hadoop.tmp.dir</name>

<value>/mnt/raid</value>

</property> <!-- Comment this property if not using the raid -->

</configuration>

2. vi **slaves**

// Add the localhost in the file, if not already added.

localhost

3. Vi **hadoop-env.sh**

// **hadoop-env.sh** has environment variable settings which is internally used by Hadoop.

// I have uncommented the java path and configured the java path.

The java implementation to use.

export JAVA_HOME=/usr/lib/jvm/java-7-openjdk-amd64

4. vi **hdfs-site.xml**

hdfs-site.xml has all the configuration or permission settings of HDFS, namenode path, and datanode paths of your local file systems.

// Made the replication factor of the data stored in HDFS as 1, to get the most performance out of the system and also to reduce the data over the instances. Data is not replicated here, as the default replication is 3 in Hadoop.

// Set the permission as false to permit all the users to use the HDFS. i.e, disabling the permission check.

// Added the namenode path and Data node path, while configuring the Multi node cluster.

<configuration>

<property>

<name>dfs.replication</name>

<value>1</value>

</property>

<property>

<name>dfs.permissions</name>

<value>>false</value>


```
</property>
</configuration>
```

5. vi **mapred-site.xml**

```
// mapred-site.xml is used to specify which MapReduce framework we
// want to use, how many mappers and reducers we have to increase.
// Set the mappers and reducers to 4 as c3.large has 2 virtual cores.
// Mappers and reducers can also be set using command line
// -D Mapred.reduce.task=4 and -D Mapred.map.task=4
```

```
<configuration>
<property>
    <name>mapred.job.tracker</name>
    <value>hdfs://ec2-54-175-61-81.compute-
1.amazonaws.com:8021</value>
</property>

<!-- Increase the per task JVM memory of mapper and reducer, if heap
error comes in between -->
<property>
    <name>mapreduce.map.java.opts</name>
    <value>-Xmx2342m</value>
</property>
<property>
    <name>mapreduce.reduce.java.opts</name>
    <value>-Xmx4684m</value>
</property>

<!-- Increase the memory of mapper and reducer, if memory not
sufficient error comes in between -->
<property>
    <name>mapreduce.map.memory.mb</name>
```

```
        <value>2928</value>
    </property>
    <property>
        <name>mapreduce.reduce.memory.mb</name>
        <value>5856</value>
    </property>

    <!-- Set the maximum limit of mapper and reducer -->
    <property>
        <name>mapred.tasktracker.map.tasks.maximum</name>
        <value>4</value>
    </property>
    <property>
        <name>mapred.tasktracker.reduce.tasks.maximum</name>
        <value>4</value>
    </property>

    <!-- Set the number of mapper and reducer to be used in the program-->
    <property>
        <name>mapred.map.tasks</name>
        <value>4</value>
    </property>
    <property>
        <name>mapred.reduce.tasks</name>
        <value>4</value>
    </property>
</configuration>

// Done with configuring the things.
<!-- OPTIONAL
// Used to turn off the safe mode of namenode.
./hdfs dfsadmin -safemode leave
```

```
// To format the namenode  
sudo bin/hadoop namenode -format  
-->
```

```
// To delete temp files  
rm -Rf /tmp/hadoop-ubuntu/
```

```
// To grant permissions  
sudo chmod 777 /mnt/raid/
```

```
// To format the namenode  
cd ../bin/  
sudo ../hdfs namenode -format
```

```
// start the dfs  
cd ../sbin/  
~/hadoop-2.7.2/sbin$ ./start-dfs.sh  
jps
```

```
// Make a director on HDFS  
cd ../bin/  
~/hadoop-2.7.2/bin$ ./hdfs dfs -mkdir -p /user/pshetty4/inputfolder
```

```
// Creating RAID on disk
```

```
sudo apt-get install mdadm  
sudo umount -l /mnt  
sudo mdadm --create --force --verbose /dev/md0 --level=0 --  
name=MY_RAID --raid-devices=3 /dev/xvdb /dev/xvdc /dev/xvdd
```

```
// Optional if mdadm says Device busy
```

```
sudo mdadm --stop /dev/md0
```

```
sudo mkfs.ext4 -L MY_RAID /dev/md0
```

```
sudo mkdir -p /mnt/raid
```

```
sudo mount LABEL=MY_RAID /mnt/raid
```

```
// Now check the mounting status using lsblk command.
```

```
lsblk
```

```
// Create 10GBUnsorted data
```

```
sudo ~/64/./gensort -a 100000000 /mnt/raid/10GBUnsorted
```

```
OR
```

```
// Without raid
```

```
sudo ~/64/./gensort -a 100000000 ~/10GBUnsorted
```

```
// PUT THE 10GBUnsorted on HDFS
```

```
~/hadoop-2.7.2/bin$ ./hadoop fs -put /mnt/raid/10GBUnsorted
```

```
/user/pshetty4/inputfolder/
```

```
OR
```

```
// Without raid
```

```
~/hadoop-2.7.2/bin$ ./hadoop fs -put ~/10GBUnsorted
```

```
/user/pshetty4/inputfolder/
```

```
//Remove the 10GB from raid locally
```

```
~/hadoop-2.7.2/bin$ rm -r /mnt/raid/10GBUnsorted
```

```
OR
```

```
// Without raid
```

```
~/hadoop-2.7.2/bin$ rm -r ~/10GBUnsorted
```

```
// Run the job on HDFS
```

```
~/hadoop-2.7.2/bin$ ./hadoop jar HadoopSortSingleNode.jar
```

```
HadoopSortSingleNode /user/pshetty4/inputfolder 10GBSorted
```


Performance Report CS-553 Spring 2016(A20354961)

```
f!6Suy2      00000000000000000000000000003ABFD84 EEEE555555556666AAAA5555BBBDDDD0000111166660000DDDD
%#NlPq.     0000000000000000000000000000003B3FB9 1111000033334444111166666666AAAAAAAAA00001111CCCEEE
%'^cl'~     0000000000000000000000000000002EDC5C8 8888AAAA11114444FFFF77773333EEEE44440000FFFF99999999
$"-QJ]      0000000000000000000000000000005F1265D CCCC6666EEEE22220000DDDAAAA88886666BBBB00006666AAAA
```

```
// Get the last ten records
```

tail -10 part-r-00000

ubuntu@ip-172-31-4-13:~/10GBSorted\$

```
ubuntu@ip-172-31-4-13:~/10GBSorted$ tail -10 part-r-00000
```

```
~~~~~uq2k#=#U      000000000000000000000000C06745 99991111DDDD22221110000FFFFFFFFFFF33337777CCCC2222
~~~~~v/O&Qnm      0000000000000000000000004709701 CCCC88883333FFF0000000000099991111FFFF777744446666
~~~~~yKOl:gE       00000000000000000000000002048B4F CCCC11114444888822226666BBBB888855557777EEEEBBBB0000
~~~~~yK^H.il       0000000000000000000000000463D004 44440000FFFF3333999944447777DDDDFFFAAAA11118888DDDD
~~~~~yL:C'XE       00000000000000000000000005B0D211 2222EEEE3333000022221111CCCFFF555577774444BBBB6666
~~~~~zbA_ Tt       000000000000000000000000007F9F4F BBBBCCCC666655559999FFFF8888AAAA11116666AAAABBBB0000
~~~~~zeO^FEg       00000000000000000000000001E06130 4444CCCB BBB99992222888855558888CCCFFF000011111111
~~~~~}GxjWHl       00000000000000000000000000CA1345 777711118888AAAAAAAA22221111BBBB00002222BBBBCCCC2222
~~~~~]P:]gOg       000000000000000000000000040DA3E4 4444FFFF444466663333EEEE88888888DDDEEEE44442222DDDD
~~~~~|ku|K<p       00000000000000000000000005E4A0AA 0000666655551111BBBB88889999AAAA55550000333355557777
```

2. Multi Node Sorting

Map and Reduce is done on the dataset on single node.

Instance type: c3.large

Type and no of slave nodes: 16 & c3.large

Type and no of slave nodes: 1 & c3.large

Hadoop version: Hadoop 2.7.2

Execution and Environment setup:

Open all instances with Hadoop installed AMI's and place the public dns and private ipaddress to configure the Multi-Node Cluster

Note: Configure one datanode initially with the configuration listed below and take a AMI in amazon Ec2 console and launch 16 more instances and list all those in slaves. This will ease the work of setting up the cluster.

Change the permissions of the pem key for SSH to data nodes.

my Linux terminal\$ **sudo chmod 600 ~/.ssh/HadoopPshetty4New.pem**

Now let's SSH into the one of the local ec2 machine with the following SSH command template:

my Linux terminal\$ **ssh -i ~/.ssh/HadoopPshetty4New.pem
ubuntu@ec2_instance_public_dns**

Setting up a config file in the ~/.ssh directory to ease of connecting to instance.

my Linux terminal\$ **~/.ssh/config** file.

//Create one, if it does not exist. Config file is also submitted along with this submission. Config includes all the ipaddress of the 1 master and 16 slave instances.

.....

// Send this config file and pem file to the namenode.

my Linux terminal\$ **scp /home/ubuntu/.ssh/HadoopPshetty4New.pem
~/.ssh/config namenode:~/.ssh**

The authenticity of host 'ec2-52-91-102-135.compute-1.amazonaws.com (172.31.4.185)' can't be established.

ECDSA key fingerprint is 30:c0:45:72:05:df:8c:3e:1f:46:14:89:44:f9:2c:7d.

Are you sure you want to continue connecting (yes/no)? yes

Warning: Permanently added 'ec2-52-91-102-135.compute-

1.amazonaws.com,172.31.4.185' (ECDSA) to the list of known hosts.

HadoopPshetty4New.pem

100% 1692 1.7KB/s 00:00

config

100% 2095 2.1KB/s 00:00

my Linux terminal\$ **ssh namenode**

// connects to the master node.

On the NameNode, create the public/private rsa keypair using
~/.ssh/id_rsa.pub,
and add it to it's authorized_keys

namenode\$ **ssh-keygen -f ~/.ssh/id_rsa -t rsa -P ""**

namenode\$ **cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys**

Copy the public/private rsa keypair to each DataNode's

~/.ssh/authorized_keys

to enable the passwordless SSH capabilities from the NameNode to any
DataNode.

namenode\$ **cat ~/.ssh/id_rsa.pub | ssh datanode1 'cat >>**

~/.ssh/authorized_keys'

namenode\$ **cat ~/.ssh/id_rsa.pub | ssh datanode2 'cat >>**

~/.ssh/authorized_keys'

.....

Try doing ssh to datanode,

namenode\$ **ssh datanode1**

.....

% NameNode and DataNode COMMON configurations %

Open /home/ubuntu/hadoop-2.7.2/etc/hadoop for all the configuration files.

For convenience, this line “/home/ubuntu/hadoop-2.7.2/etc/hadoop” has been used as \$HADOOP_CONF_DIR.

\$HADOOP_CONF_DIR :- /home/ubuntu/hadoop-2.7.2/etc/hadoop

sudo vim \$HADOOP_CONF_DIR/hadoop-env.sh

// Uncomment and add this below line

The java implementation to use.

export JAVA_HOME=/usr/lib/jvm/java-7-openjdk-amd64

sudo vim \$HADOOP_CONF_DIR/core-site.xml

// change the localhost to namenode_public_dns in all datanodes

<configuration>

<property>

<name>fs.defaultFS</name>

<value>hdfs://namenode_public_dns:9000</value>

</property>

</configuration>

sudo vim \$HADOOP_CONF_DIR/yarn-site.xml

// add these configuration lines in it

<configuration>

<!-- Site specific YARN configuration properties -->

<property>

```
<name>yarn.nodemanager.aux-services</name>
<value>mapreduce_shuffle</value>
</property>
<property>
  <name>yarn.nodemanager.aux-
services.mapreduce.shuffle.class</name>
  <value>org.apache.hadoop.mapred.ShuffleHandler</value>
</property>
<property>
  <name>yarn.resourcemanager.hostname</name>
  <value>namenode_public_dns</value>
</property>
</configuration>
```

```
// copy the mapred-site.xml from the template
sudo cp $HADOOP_CONF_DIR/mapred-site.xml.template
$HADOOP_CONF_DIR/mapred-site.xml
sudo vim mapred-site.xml
```

```
<configuration>
  <property>
    <name>mapreduce.jobtracker.address</name>
    <value>namenode_public_dns:54311</value>
  </property>
  <property>
    <name>mapreduce.framework.name</name>
    <value>yarn</value>
  </property>
</configuration>
```

% NameNode specific configurations %

// adding all the hosts ipaddresses to /etc/hosts

sudo vi /etc/hosts

ec2-52-91-102-135.compute-1.amazonaws.com ip-172.31.4.185

.

.

ec2-52-87-151-144.compute-1.amazonaws.com ip-172.31.1.99

// Open up the hdfs-site.xml

sudo vim \$HADOOP_CONF_DIR/hdfs-site.xml

<property>

<name>dfs.namenode.name.dir</name>

<value>file:///home/ubuntu/hadoop2.7.2/hadoop_data/hdfs/namenode

</value>

</property>

Add a masters file to config directory and add namenode hostname to it.

sudo touch \$HADOOP_CONF_DIR/masters

sudo vim \$HADOOP_CONF_DIR/masters

Open up the slave file and add all the datanodes hostname to it.

sudo vim \$HADOOP_CONF_DIR/slaves

.....

%% Data Node configurations %%

// Open up the hdfs-site.xml

```
sudo vim $HADOOP_CONF_DIR/hdfs-site.xml
```

```
<property>
<name>dfs.namenode.name.dir</name>
<value>file:///home/ubuntu/hadoop2.7.2/hadoop\_data/hdfs/datanode
</value>
</property>
```

.....

%% Starting the cluster %%

Format the name node first and then launch all dfs and yarn, either separately using **start-dfs.sh** and **start-yarn.sh**

```
namenode$ ~/hadoop-2.7.2/bin/hdfs namenode -format
namenode$ ~/hadoop-2.7.2/sbin/start-dfs.sh
namenode$ ~/hadoop-2.7.2/sbin/start-yarn.sh
```

OR

```
namenode$ ~/hadoop-2.7.2/bin/hdfs namenode -format
namenode$ ~/hadoop-2.7.2/sbin/start-all.sh
```

```
namenode$ ~/hadoop-2.7.2/sbin/jps
```

8176 SecondaryNameNode

8623 Jps

7910 NameNode

8334 ResourceManager

Try in any of the datanode whether it is connected or not.

```
datanode$ jps
```

1580 NodeManager

1464 DataNode

1724 Jps

.....

%% Running the application %%

// Generate 100GB unsorted data using gensort
sudo ~/64/./gensort -a 1000 ~/hadoop-2.7.2/bin/100GBUnsorted

// make your own directory in HDFS
./hdfs dfs -mkdir -p /user/pshetty4/inputfolder

// Put the data in HDFS
./hadoop fs -put 100GBUnsorted /user/pshetty4/inputfolder/

// Remove the local 100GBUnsorted file once uploaded to HDFS
rm 100GBUnsorted

// Run the jar
**./hadoop jar HadoopSortSingleNode.jar HadoopSortSingleNode
/user/pshetty4/inputfolder 100GBSorted**

// Check the hdfs for file
./hadoop fs -ls

// get the sorted file from hdfs locally
./hadoop fs -get 100GBSorted .

// Traverse till the output file and convert the file to dos to run valsart
**cd 100GBSorted/
unix2dos part-r-00000
~/64/./valsart part-r-00000**

```
ubuntu@ip-172-31-4-185:~/hadoop-2.7.2/bin/100GBSorted$  
~/64/./valsort part-r-00000
```

```
// Output
```

Records: 1000000000

Checksum: 1dcd7024ab4bf409

Duplicate keys: 0

SUCCESS - all records are in order

```
// Get the top ten sorted keys
```

head -10 part-r-00000

```
ubuntu@ip-172-31-4-185:~/hadoop-2.7.2/bin/100GBSorted$ head -10
part-r-000000
```

[illegible]

```
// Get the last ten sorted keys
```

tail -10 part-r-00000

```
ubuntu@ip-172-31-4-185:~/hadoop-2.7.2/bin/100GBSorted$ tail -10
part-r-000000
```

```
~~~~~#iyaiX      000000000000000000000000025D35EDF   6666AAAA5555999977770000222333888FFFF99992220000  
~~~~~+@p){@      0000000000000000000000000085426F4   777733335555111111110000CCCC55559999AAA7777DDDDDDDD  
~~~~~,R'_?n     00000000000000000000000001034E347   111111119999000011118888AAAA5555444EEEF999933338888  
~~~~~,Ey^')     000000000000000000000000016FO66EB   CCCC6666DDDD2222DDDD111188889999EEEEEEEEEEEBBBBB4444  
~~~~~4!kA7x     00000000000000000000000001F1A1E26   EEEF777711117777BBB1111EEEE88884444DDDDDDDEE EEBBBB  
~~~~~8li!/|@    00000000000000000000000001F05932F   11119999BBBB44447777000011114444CCCAAAA6666DDDD0000  
~~~~~<'!5>F     000000000000000000000000008CB2293   88883333BBB1111666699998888555588888882228888CCCC  
~~~~~(G-)m^!)   000000000000000000000000013397f73         DDDDFFFB BBBCCCCCFFF44446666AAA1111333333AAACCCC  
~~~~~c+l&cP     000000000000000000000000074BDf64        8888000055550000DDDD2227777AAA00003333222AAAADDD  
~~~~~hb&$*X     0000000000000000000000000032CF0E6B       7777BBBBBBBB9999EEEEAAAAA0000CCCCDD DD4444BBBB4444
```


Implementation:

- Hadoop provides **Mapper** and **Reducer** interfaces and the same has been implemented as MapReduce program in java. It provides map and reduce methods. In map, it transforms the input records into intermediate splits and the output of this phase is given to input of reducer phase. In reducer phase, it reduces the splits and the data is sorted with sort and shuffle phases and the output is written to the file system.
- Followed the **environmental setup** that has been explained above. (Launching the instances, configuring it to run, making executable jar file of sort functionality)
- Created the 10GB dataset on single node and 100GB dataset on Multi node mode using **Gensort** and put it in the HDFS and ran the jar to sort the data across all the nodes.
- Once the experiment is completed, we can get the output file from the HDFS locally and perform **Valsort** to check the sort correctness.

Note: The experiment has been done again by changing the number of mappers and reducers to 4 using **mapred-site.xml** in conf directory **mapred.map.tasks** and **mapred.reduce.tasks**. Decrease in time taken to process the sorting was noticed when mappers and reducers were increased.


Execution Chart of Hadoop on 10GB and 100GB Dataset:

Data-size	Execution Time(in secs)
10GB	1078
100GB	10286



Throughput of Hadoop:

Data-size	Throughput(in MBps)
10GB	9.276437848
100GB	9.721952168



HADOOP Q & A's

1) What is a Master node? What is a Slaves node?

Ans: Name node and the secondary name node in Hadoop architecture are termed as Master nodes.

Data nodes are termed as Slaves node.

- The **Name node**, the central controller of HDFS, holds all the HDFS **metadata** for the multi node cluster. It knows only what blocks makes up a file and where those blocks are located in HDFS. It has **ResourceManager** service, which is the heart of YARN, checks the scheduling of application tasks and management of the Hadoop cluster's resources.
- The **Data node**, are where Hadoop data is actually stored and where data processing takes place. It has **NodeManager** service, coordinates the resources for an individual slave node and reports back to the Resource Manager.

2) Why do we need to set unique available ports to those configuration files on a shared environment? What errors or side-effects will show if we use same port number for each user?

Ans: Hadoop uses a lot of ports for its internal and external communications. Usually, you can only have one application (or services) listening on a single port at one time, that's the reason why ports exist, so that it allows multiple services in Hadoop architecture to share the network without any conflicts.

Errors/Side effects:

The port is in use OR there is an instance of the service already running OR Address already in use OR Ports are not available.

3) How can we change the number of mappers and reducers from the configuration file?

We can change it using **mapred-site.xml** which is available in conf directory of Hadoop,

```
<property>
  <name>mapred.map.tasks</name>
  <value>4</value>
</property>
<property>
  <name>mapred.reduce.tasks</name>
  <value>4</value>
</property>
```

// This can be set also in command prompt using "-D **Mapred.reduce.task=4**" and "-D **Mapred.map.task=4**"

(iii) Spark sort:

The Apache Spark software library is a fast engine that allows for the distributed processing of large data sets across clusters of computers. It can run the programs up to 100x faster than Hadoop. It runs on top of Hadoop also and can access many data sources like HDFS, Cassandra, etc. It provides a powerful tool, **Spark Shell** (which runs on JVM) to analyze data and can run Scala, Python languages. It uses Hadoop's HDFS and YARN to run jobs on big data. **Resilient Distributed Datasets (RDDs)** is created by referencing a dataset in HDFS.

1. Single Node Sorting

Instance type: c3.large

Hadoop version: Spark 1.6.0

Execution and Environment setup:

** Copy the Spark folder into the instance,

scp -i spark_node.pem /spark_folder ubuntu@spark_node_public_dns

**** connect to the instance,**
ssh spark_node_public_dns

**** Add the pem file**
eval `ssh-agent -s`
chmod 600 SparkPshetty4pem.pem
ssh-add SparkPshetty4pem

**** Copy the **SortTheKeySingleNode.py** file to the bin.**

**** Using gensort, create 10GB of unsorted data,**

./gensort -a 100000000 10GBUnsorted

**** Submit the python file using spark-submit to run the job.**
ubuntu@ip-172-31-3-205:~/spark-1.6.0-bin-hadoop2.6/bin\$ **./spark-**
submit SortTheKeySingleNode.py

**** Once the job is completed, traverse till the output file convert it to dos format.**

ubuntu@ip-172-31-3-205:~/spark-1.6.0-bin-hadoop2.6/bin\$ **cd**
10GBSorted/
ubuntu@ip-172-31-3-205:~/spark-1.6.0-bin-hadoop2.6/bin/10GBSorted\$
unix2dos part-00000
ubuntu@ip-172-31-3-205:~/spark-1.6.0-bin-hadoop2.6/bin/10GBSorted\$
~/64/valsort part-00000

// Output

Records: 359637

Checksum: 2bf1116d3c69f

Duplicate keys: 0

SUCCESS - all records are in order


```
~~~}GxjWHI 0000000000000000000000000000CA1345 777711118888AAAAAAA22221111BBBB00002222BBBBCCCC22 22
~~~}P;]gOg 000000000000000000000000000040DA3E4 4444FFFF444466663333EEEE88888888DDDDDEEEE44442222DDDD
~~~}kU|K<p 000000000000000000000000000005E4A0AA 0000666655551111BBBB88889999AAAA55550000333355557777
```

2. Multi Node Sorting

Instance type: c3.large

Hadoop version: Spark 1.6.0

Execution and Environment setup:

*** Once single node is done, create a AMI of it without volume and open up a new instance with this AMI as it already has spark installed and open the ec2 folder of spark directory. Copy the pem file which should be used to launch cluster.

```
localTerminal$ scp -i SparkPshetty416Nodes.pem
SparkPshetty416Nodes.pem ubuntu@ip-172-31-49-133/~
```

```
ubuntu@ip-172-31-49-133:~/spark-1.6.0-bin-hadoop2.6$ cd spark/ec2
```

*** Export the AWS_ACCESS_KEY_ID and AWS_SECRET_ACCESS_KEY to the ec2 folder.

```
ubuntu@ip-172-31-49-133:~/spark-1.6.0-bin-hadoop2.6/ec2$ export
AWS_ACCESS_KEY_ID=*****
```

```
ubuntu@ip-172-31-49-133:~/spark-1.6.0-bin-hadoop2.6/ec2$ export
AWS_SECRET_ACCESS_KEY=*****
```

*** Change the mode for SSH to work between master and workers

```
ubuntu@ip-172-31-49-133:~/spark-1.6.0-bin-hadoop2.6/ec2$ chmod 400
SparkPshetty416Nodes.pem
```

*** Using the AMI which has spark and hadoop installed, Launch a 16 nodes cluster using "spark-ec2" command.

```
***** ./spark-ec2 -k <keypair>
        -i <key-file> -s <num-slaves> --ami <ami-id>
        launch <cluster-name>,
        where <keypair> is the name of your EC2 key pair,
        <key-file> is the private key file for your key pair,
        <num-slaves> is the number of slave nodes to launch
        <cluster-name> is the name to give to your cluster
<ami-id> name generated in Amazon EC2 webconsole. *****
```

```
ubuntu@ip-172-31-49-133:~/spark-1.6.0-bin-hadoop2.6/ec2$ ./spark-
ec2 -k SparkPshetty416Nodes -i SparkPshetty416Nodes.pem -s 16 --
instance-type=c3.large --region=us-east-1 --ami= ami-940ee4f4 --spot-
price=0.035 --hadoop-major-version=yarn launch 16NodesSpark100GB
```

.....

*** Now connect to the master node using the pem file and with root as an user, as it is amazon AMI.

Before this copy the gensort and valsort and python sort file to the master instance.

```
My Linux Terminal$ scp -i ~/SparkPshetty416Nodes.pem -r 64/
root@ec2-54-200-18-218.compute-1.amazonaws.com:~
My Linux Terminal$ ssh -i ~/SparkPshetty416Nodes.pem root@ ec2-54-
200-18-218.compute-1.amazonaws.com
```

*** Create the 100 GB unsorted file using gensort

```
root@ip-172-31-31-55$ ~/64/gensort -a 1000000000 100GBUnsorted
```

*** Go to hadoop directory and put the unsorted file into HDFS.

[illegible]

```
// Get the last ten records
```

tail -10 part-r-00744

```
~~~~~#iaylX 00000000000000000000000025D35EDF 6666AAAA55559999777700002223338888FFFF999922220000
~~~~~+@p|{@ 0000000000000000000000000085426F4 777733335555111111110000CCCC55559999AAAA7777DDDDDDDD
~~~~~,R^_?n 000000000000000000000000001034E347 111111119999000011118888AAAA55554444EEEE999933338888
~~~~~,Ey^A) 0000000000000000000000000016F0E66B CCCC6666DDDD2222DDDD111188889999EEEEEEEEEEEBBBB4444
~~~~~4!kA7x 000000000000000000000000001F1A1E26 EEEE777711117777BBBB1111EEEE88884444DDDDDDDDDEEEEBBBB
~~~~~8li!/ @ 000000000000000000000000001F05932F 11119999BBBB44447777000011114444CCCCAAA6666DDDD0000
~~~~~<'5>F 0000000000000000000000000008CB2293 88883333BBBB111166669999888855558888888822228888CCCC
~~~~~(G-)m^!) 0000000000000000000000000013397F73 DDDDFFFFBBBBCCFFFF44446666AAAA111133333333AAACCCC
~~~~~c+I&cP 00000000000000000000000000074BDf64 8888000055550000DDDD22227777AAAA000033332222AAADDDD
~~~~~hb&5*X* 0000000000000000000000000032C0E06B 7777BBBBBBBBBB9999EEEEAAAAAAA0000CCCCDDDD4444BBBB4444
```

*** TO stop the cluster after the job is run use destroy.

```
./spark-ec2 destroy 16NodesSpark100GB
```


Implementation:

- Python library has been used to implement.
- Followed the environmental setup that has been explained above. (Launching the instances/Clusters with proper Disk size, connecting it to master node)
- Created the 10GB dataset on single node and 100GB dataset on Multi node mode using Gensort and put it in the HDFS and ran the "SortTheKeySingleNode" python file to sort the data across all the nodes. The input is divided into individual key value RDD pairs with first 10 bytes as key and the next 90 bytes as value. Used **map()** do the same. And then with this result, used **sortBykey()** function to sort the data and using **saveAsTextFile()** function stored it to the HDFS.
- Once the experiment is completed, we can get the output file from the HDFS locally and perform Valsort to check the sort correctness.
- I have used c3.large instance for this and made major changes of SPARK_EXECUTOR_INSTANCES to 4, SPARK_EXECUTOR_MEMORY to 2GB, SPARK WORKER MEMORY to 2,

SPARK_EXECUTOR_INSTANCES to 4 in **spark-env.sh**, to get the best performance out of the run.


Execution Chart:

Data-size	Execution Time(in secs)
10GB	2007
100GB	2435



Throughput Chart:

Data-size	Throughput(in MBps)
10GB	4.982561036
100GB	41.0677



Performance related Observations

1. Compare the performance of the three versions of Sort (Shared-Memory, Hadoop, and Spark) on 1 node scale and explain your observations. Compare the Shared-Memory performance of 1 node to Hadoop and Spark Sort at 16 node scales and explain your observations.

Answer:


Time taken chart of all 3 sort(in secs):

Time taken to sort in secs



Data-Size	Shared Memory	Hadoop	Spark
10GB	422	1078	2007
100GB	NA	10286	2435

Speedup chart of Hadoop and Spark with respect to Shared Memory, 10GB Dataset:

FrameWork	SpeedUp with respect to Shared Memory
Shared Memory	1
Hadoop	0.410018553
Spark	0.210264076



Throughput comparison among all 3 sort versions for 10 GB and 100GB:



From the above graphs and the value, we can infer that SharedMemory using External Sort performs well compared to Hadoop and Spark. But when taken into consideration about handling big data, Hadoop performs better on single node than Spark.


2. Draw an execution line chart and a speed up line chart for 1 node and 16 node cases, for Java, Hadoop, and Spark Sort. For Spark and Hadoop, compute two different speedups, using different base cases; one speedup (speedup-shared-memory) should be relative to the Shared-memory Sort performance (note that you might get a speedup less than 1); the second speedup (speedup-spark & speedup-hadoop) should be relative to the Spark or Hadoop performance at 1 node scale respectively (should be a number greater than 1). **What conclusions can you draw?** Which seems to be **best at 1 node scale**? How about **16 nodes**? Can you predict which would be best at **100 node scale**? How about **1000 node scales**?

Answer:

Speedup chart of Spark with respect to Hadoop, 100GB Dataset:

Hadoop in secs	Spark in secs
10286	2435

FrameWork	SpeedUp of Spark with respect to Hadoop
Hadoop	1
Spark	4.224229979



Best at one node scale: Hadoop

Best at 16 node scale: Scala

Best at 100 node scale: Scala

Best at 100 node scale: Scala

References:

<http://www.dummies.com/how-to/content/master-nodes-in-hadoop-clusters.html>

<http://www.dummies.com/how-to/content/slave-nodes-in-hadoop-clusters.html>

<http://bradhedlund.com/2011/09/10/understanding-hadoop-clusters-and-the-network/>

<http://spark.apache.org/docs/latest/programming-guide.html#resilient-distributed-datasets-rdds>

<https://asadrazajaffree.wordpress.com/2014/05/14/multi-node-hadoop-cluster/>

<https://cloudcelebrity.wordpress.com/2013/08/14/12-key-steps-to-keep-your-hadoop-cluster-running-strong-and-performing-optimum/>

<http://www.ashishsharma.me/2011/08/external-merge-sort.html>

<http://lemire.me/blog/2010/04/06/external-memory-sorting-in-java-the-first-release/>

<http://spark.apache.org/docs/latest/index.html>

https://hadoop.apache.org/docs/r1.2.1/single_node_setup.html

<http://spark.apache.org/docs/latest/programming-guide.html>

<http://hadoop.apache.org/>